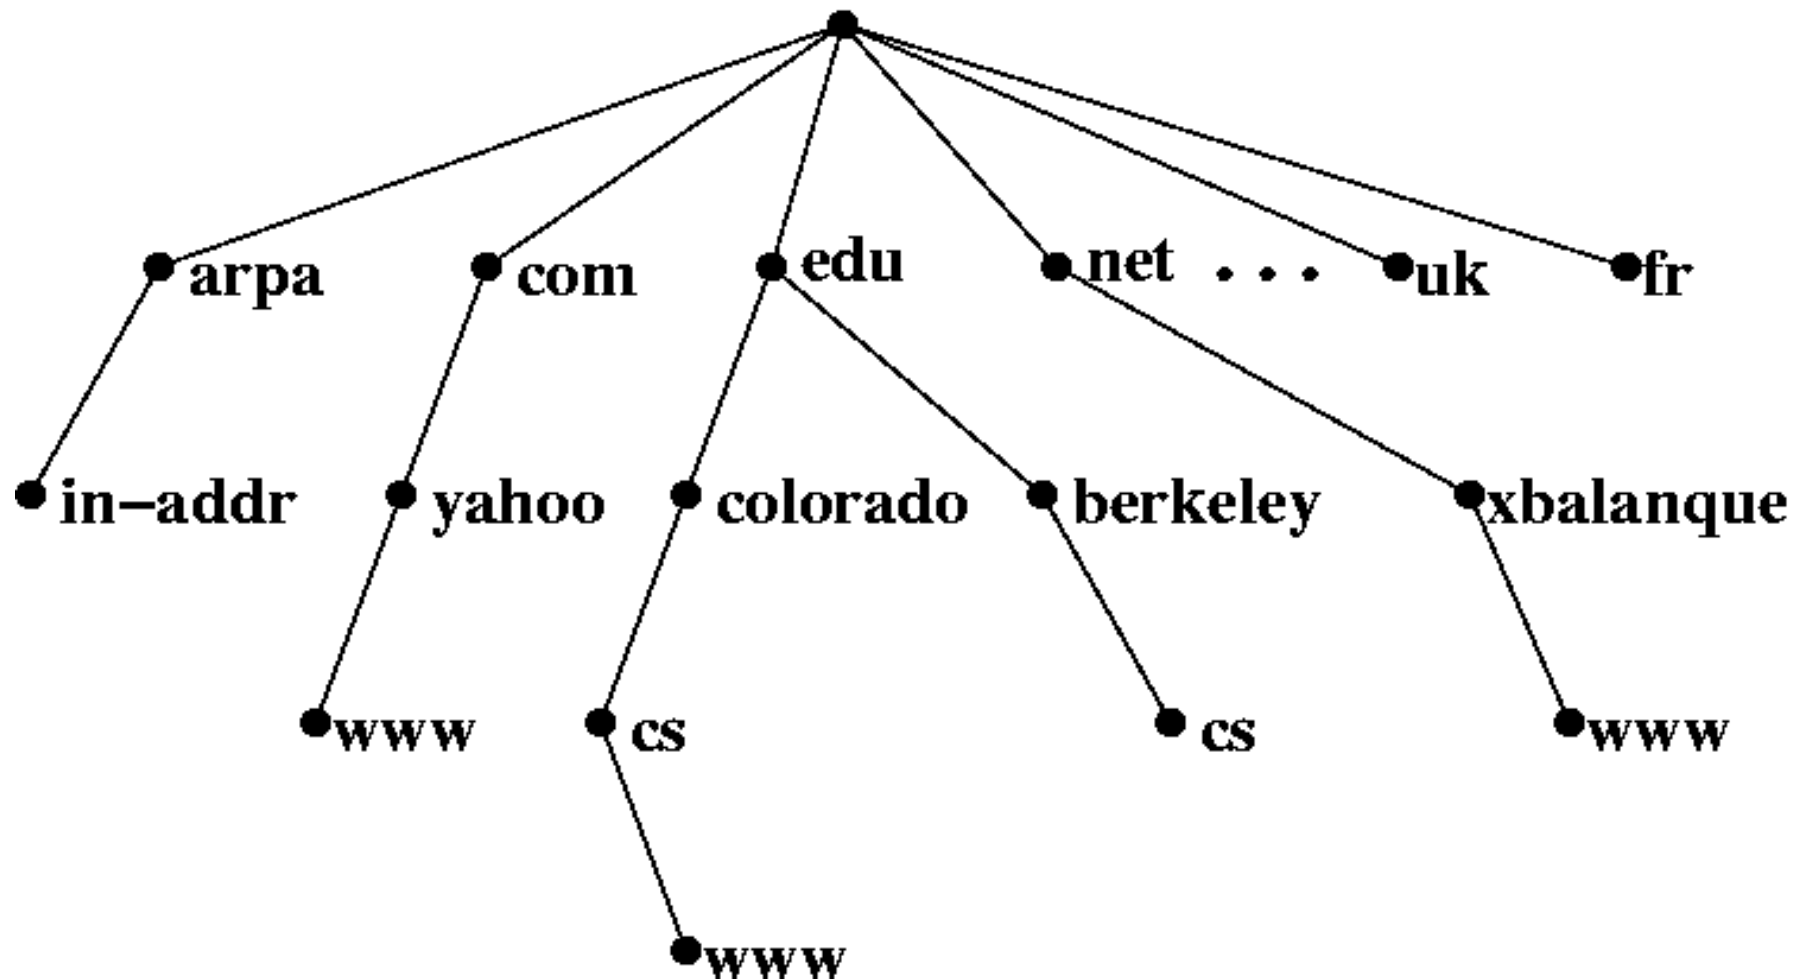


# DNS and CDNs



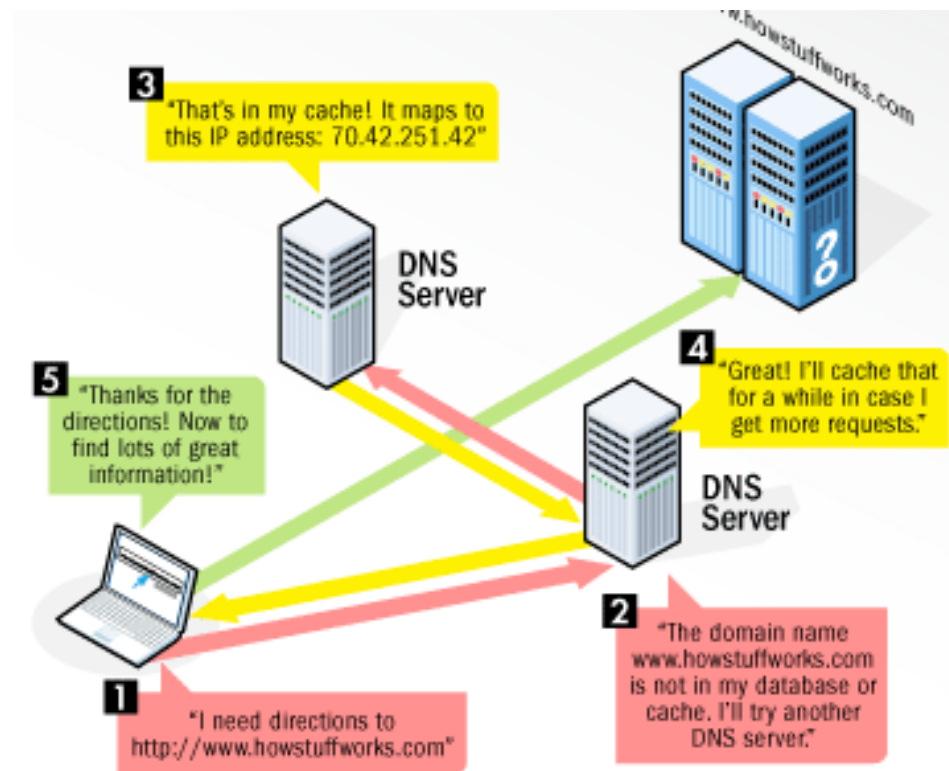
# Locating Things Online

- **Addresses** describe a location
  - A map is nice, but not for lookup
  - IP addresses, street addresses
- **Names** are mapped to addresses
  - Domain names; PO Boxes; email addresses
- **Content-based names** use the actual content to find the destination
  - Basis of publish/subscribe and peer-to-peer

# Domain Name System

- The Internet's routers know only IP addresses
  - Fine for routers, not great for humans
- DNS translates domain names to IP addresses
  - `google.com => 74.125.95.99`
  - (or something similar)
  - The "phone book" of the internet

# DNS at work




# DNS Design

- DNS is a globally-accessible database of name/IP pairs. Some requirements:
  - Needs to be up all the time
  - Continuously updated by many parties
  - Must be accurate; errors prevent connections
  - Serves massive query load
  - Needs distributed administration
  - Source of political & commercial disputes
  - Potential terrorism target
- Overall: utterly shocking that it works

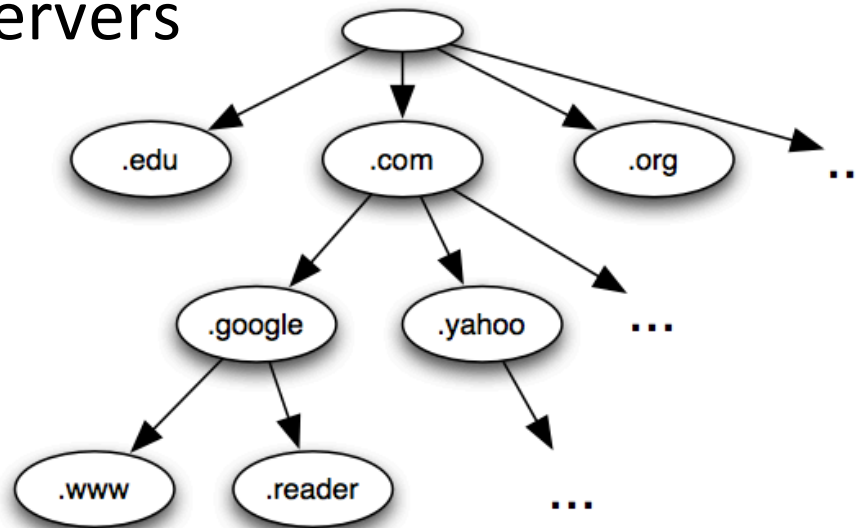
# The Namespace

- Tree structure, 1-63 chars per node
- *Fully-qualified domain name* is leaf-to-root name
- Only *DNS root* has no parent




# The Namespace

- Nodes grouped into administrative zones
  - E.g., root, com, google.com
- Each zone served by authority servers
  - AKA authoritative name servers
- Authority server can delegate subdomains to other authority servers



# The Namespace

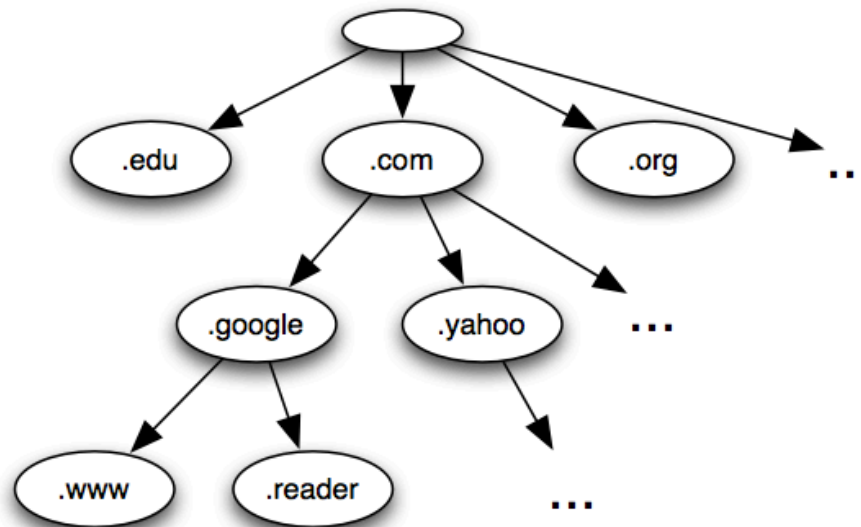
- Servers are primary or secondary
  - Primary authority servers are given content by admins
  - Secondary authority servers grab from primaries
- A domain registrar inserts your name into the primary authority server for .com (or .net, .org, etc)





# The Namespace

- Purchased domain names inserted into 1 primary, 1 secondary (in case of primary failure)



# Try it yourself

- Check if a domain name is registered

```
$ whois umich.edu
```

```
...
```

```
Registrant:
```

```
University of Michigan -- ITD
```

```
ITCS, Arbor Lakes
```

```
4251 Plymouth Road
```

```
Ann Arbor, MI 48105-2785
```

```
UNITED STATES
```

```
...
```

```
Domain record activated:      07-Oct-1985
```

```
Domain record last updated: 14-Feb-2014
```

```
Domain expires:              31-Jul-2016
```

# The Data

- Content of DNS consists of *resource records*
  - Host/IP and IP/Host are most popular type
  - Also: zone info, email servers, other info
- Most DNS activity is this:
  - DNS client sends request to authority servers, receives resource record in response
- DNS clients built into network libraries
  - Clients use UDP (not TCP) to grab data
  - If your connection is taking a long time to establish, possibly waiting for DNS

# Resolution

- Single authority server may not be enough for a frequently queried domain name

# Try it yourself

- Look up IP address for web.eecs.umich.edu

```
$ nslookup web.eecs.umich.edu
```

```
Server:          192.168.1.1
```

```
Address:         192.168.1.1#53
```

```
Non-authoritative answer:
```

```
Name:   web.eecs.umich.edu
```

```
Address: 141.212.113.110
```

- Non-authoritative means it's returning cached results

# Try it yourself

- Look up IP address for `eeecs.umich.edu`

```
$ nslookup -type=ns eeecs.umich.edu
```

```
Server:          192.168.1.1
```

```
Address: 192.168.1.1#53
```

Non-authoritative answer:

```
eeecs.umich.edu    nameserver = csedns.eeecs.umich.edu.
```

```
eeecs.umich.edu    nameserver = dns.eeecs.umich.edu.
```

```
eeecs.umich.edu    nameserver = eeecsdns.eeecs.umich.edu.
```

Authoritative answers can be found from:

```
eeecsdns.eeecs.umich.edu internet address = 141.213.4.4
```

```
csedns.eeecs.umich.edu  internet address = 141.212.113.4
```

- Return authoritative name servers (complete zone information)

# Try it yourself

## • Alternative: dig

```
$ dig eecs.umich.edu
; <<>> DiG 9.8.3-P1 <<>> eecs.umich.edu
;; global options: +cmd
;; Got answer:
;; ->>HEADER<- opcode: QUERY, status: NOERROR, id: 29174
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 0

;; QUESTION SECTION:
;eecs.umich.edu.                IN      A

;; ANSWER SECTION:
eecs.umich.edu.                900     IN      A      141.212.113.199

;; Query time: 35 msec
;; SERVER: 192.168.1.1#53(192.168.1.1)
;; WHEN: Mon Nov 23 09:14:48 2015
;; MSG SIZE rcvd: 48
```

# Reverse DNS lookup

- Reverse DNS lookup
- Start with IP, find name

```
$ dig -x 141.212.113.199
```

OR

```
$ dig +short -x 141.212.113.199  
www.eecs.umich.edu
```



# Exercise

- What domain hosts this website?
  - `http://andrewdeorio.com`
- HINT: use both a DNS lookup and a reverse lookup
- Extra credit: Linux one-liner

# Exercise

- Where is this website hosted?
- andrewdeorio.com
- DNS lookup

```
$ dig +short andrewdeorio.com  
141.212.113.110
```

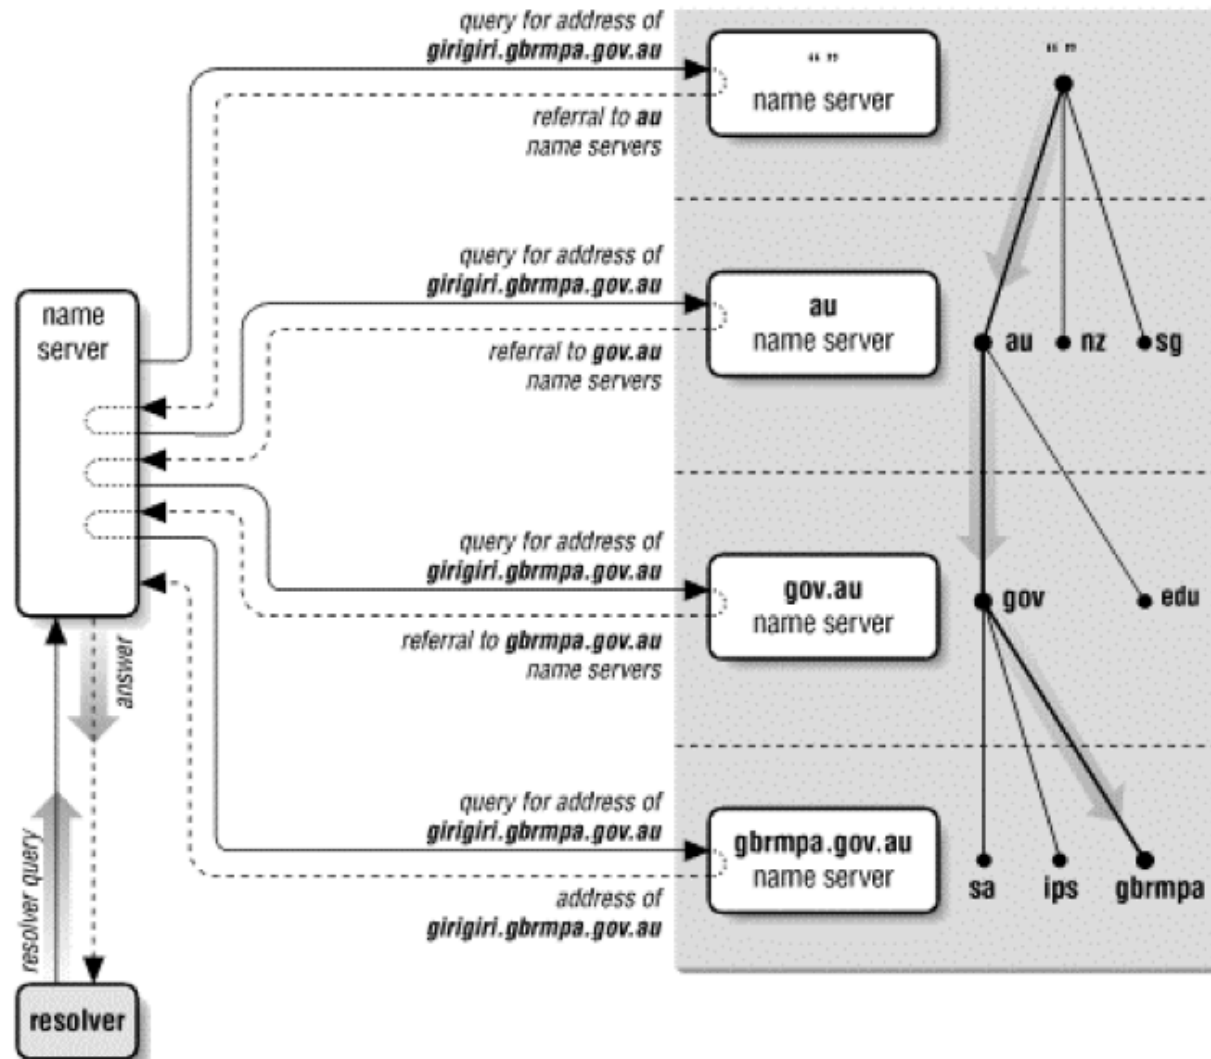
- Reverse DNS lookup

```
$ dig +short -x 141.212.113.110  
web.eecs.umich.edu.
```

- One-liner

```
$ dig +short andrewdeorio.com | xargs dig +short -x  
web.eecs.umich.edu.
```

# Resolution



# Caching

- There may be a chain of caching DNS servers between client and authority server
  - There's one for CSE
  - Caches DNS requests; answers new requests from cache whenever possible
- Great! But when data changes?
  - Each resource record has time-to-live (TTL) in seconds
  - Counts down from moment authority server emits resource record (RR)
  - Caches and clients must throw out resource records with expired time-to-live

# Root Nameservers

- Responsible for locating the top level domain name servers (.com, .net, .org, ...)
- Thirteen root name servers in world
  - Their locations are hard-coded in resolving DNS servers
  - Due to caching, involved in few queries

# Administration

- Until 1999, all US top level domains run by IANA
  - (At the time, IANA = Jon Postel)
  - IANA: Internet Assigned Numbers Authority
- Now, non-profit ICANN administers set of for-profit registrars (under contract with government)
  - ICANN: Internet Corporation for Assigned Names and Numbers

# Administration

- For a time, looked like ICANN might be transferred to United Nations. Probably won't be
- Postel's "Law": Be liberal in what you accept, and conservative in what you send

# DNS Challenges

- Shocking scale
  - Every device, service, etc
  - Spam-filtering email requires 10+ DNS lookups per message
- Extreme security vulnerability
  - What if someone can remap **google.com**?
  - What about just denial-of-service?



# DNS Cache Poisoning

- Bad info is inserted into a DNS server and cached
  - Both inadvertent and malicious
  - How could one attack CSE's DNS server?
- Remedies
  - DNSSEC requires that DNS entries be cryptographically signed
    - Signed by whom?
      - Chain of Trust
  - If you use HTTPS/SSL/TSL, problem mitigated

# DNS Cache Poisoning

- DNS cache poisoning is one of the ways that China's Great Firewall works
- Just direct requests to twitter.com to the wrong IP



# Invalid Domain Name

- DNS returns `NXDOMAIN`
- Browser shows user an error page
- Wasted “eyeball” opportunity
- Sell advertising instead
  - By inserting default entries into DNS
- Can use same trick to steal cookies
  - Insert DNS entry for non-existent sub-domain and get user to request it somehow

# Content Distribution

- Now, on to Content Distribution Networks (CDNs)

# Proxy

- Also called "proxy server"
- Looks like a server to the client, looks like client to the server
- Can have any number of proxies in the middle between client and server
- Proxy can
  - Change request or not
  - Announce itself or not
  - Change response or not
- Used for a variety of functions

# Proxy Uses


- Anonymity
- Content Filtering
- Cache Sharing

# Akamai

- Stores images, videos, at many locations throughout the world
  - Large read-only data stored close to client
  - Reduce latency to faraway datacenter
  - Reduce bandwidth costs by sending data only a short distance
- The 1st CDN (content-delivery network)
  - You might think the image is from Yahoo, but Akamai is serving it
  - Yahoo pays Akamai to do so
- Built on top of giant DNS hack!

# How it works: step 1


- Akamai places servers across country






## How it works: step 2

- Akamai gets client data (videos, imgs), copies to all nodes in network




## How it works: step 3

- How does client find nearby data?
  - Network conditions can change




# How it works: step 3

- Client rewrites URLs, uses new URLs
  - `http://yahoo.com/...` =>  
`http://akadns.akamai.net/....`




# How it works: step 3

- URL contains Akamai-controlled name
  - Specialized Akamai DNS server resolves name to nearby server; tiny TTL



# How it works: step 3

- Resolution strongly dependent on
  - Location of client; network conditions; load on Akamai servers; traffic estimation errors



# Akamai Recap

- Uses custom DNS servers to dynamically direct clients to right caching servers

# Push Technologies

- Email is "push"
  - Sender proactively sends information to the recipient
  - Good (commercial) senders require explicit subscriptions to send
- Web is "pull" only
  - HTTP is asymmetric request/response
  - Web Server cannot push content to user with an explicit user request
  - Address by having the user subscribe

# RSS

- RSS = Rich Site Summary, or Really Simple Syndication
- Pack updates to your site into an XML file using RSS tags
- Consider this a broadcast "channel"
- RSS-reader periodically checks specified sites for updates, and automatically downloads them



# Atom

- Do exactly the same steps as described for RSS, except with a slightly different specification
  - Differences are minor and in the details
  - Atom is newer and "better"
- Use Atom Syndication Format to define a "feed" as a set of "entries" in XML
- Use Atom Publishing Protocol to GET (and also to POST) the feed XML file

# Website Syndication

- Content creator places content in files created according to Atom or RSS
- Content duplicator accesses syndicated content (headlines) and displays on their own website.
- Very popular model for Blogs
- Podcasts are similar

# Writing to the Server

- Primary focus has been reading from the website, and writing back very limited information, e.g. in forms
- How about uploading data sets, videos, large formatted blog entries, ...
- SWORD: A Simple Web-Service Offering Repository Deposit
  - Based on Atom
  - Supported by many repositories

# Publish Subscribe

- In email, or other ordinary messaging, the sender identifies the recipient(s)
- In broadcast, the sender just publishes, and whoever wants to can tune in
- In subscription, the recipient specified precisely what messages it is interested in receiving
- Publish/subscribe extends the traditional messaging model in that the sender may not know who the recipients will be when creating message

# Implementation Models

- The basic concept is the same, but there are several implementations:
- List-based
  - Subscriber adds themselves to a list. Publisher sends to whoever is on the list at the time of message sending
- Broadcast
  - Publisher sends to all; local filters select
- Content-based
- Distribution Networks