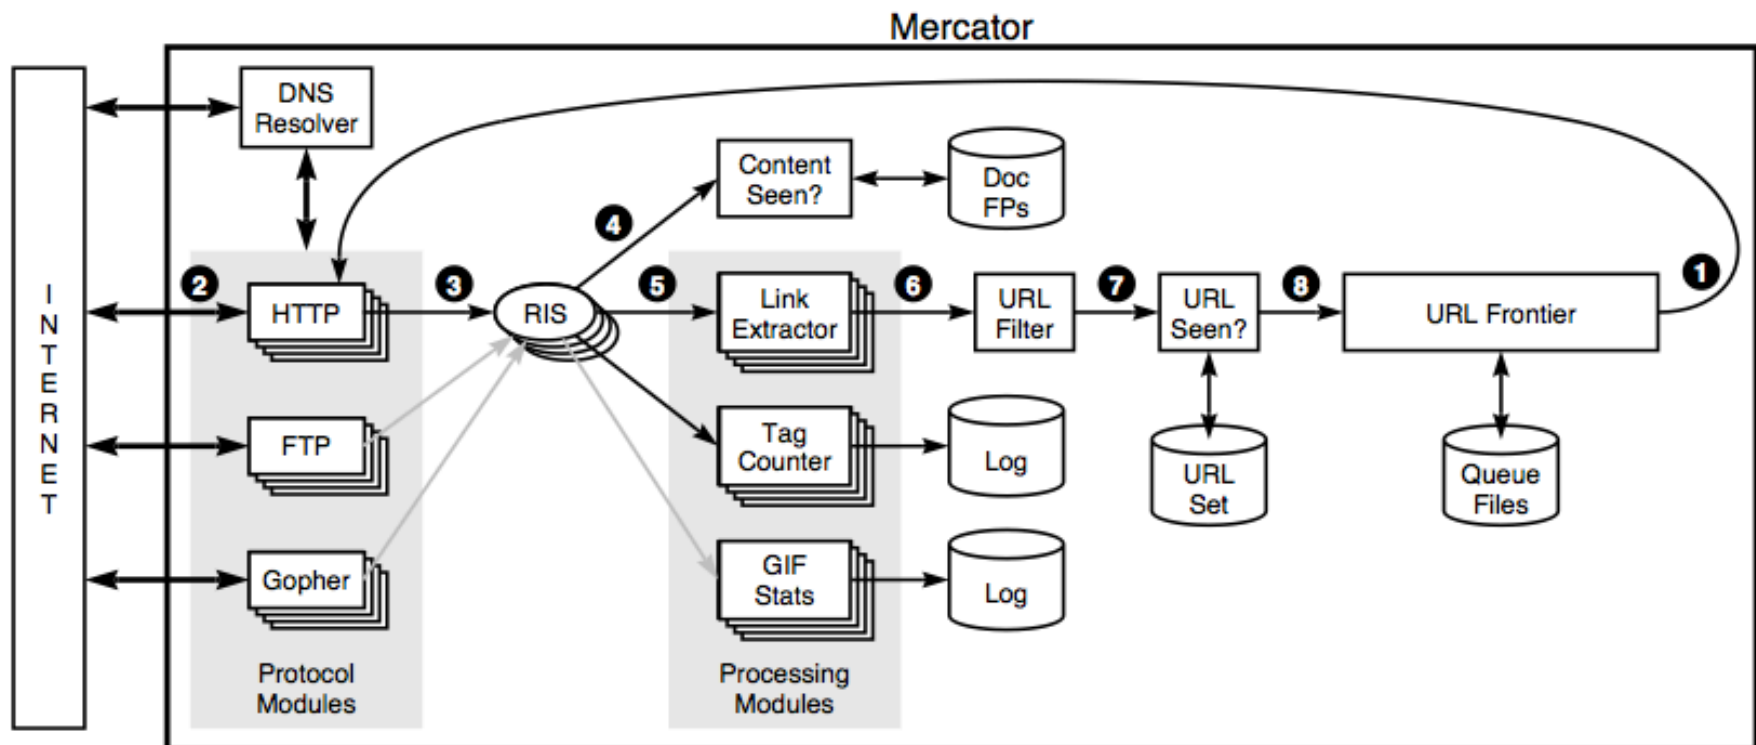


# IR3: Web Search Implementation



# Challenges

- Three challenges in web search:
  - Result relevance
  - Processing speed
  - Scaling to many documents
- So far we've discussed result relevance
- Today we'll cover speed and scaling

# A Few Numbers

- 5 B - 100 B pages
  - Let's say 10 Billion for concreteness
- Assume 10KB per compressed page
  - 100 TB data to index
- 1 minute-1 month freshness
- 3-5B queries *per day* on Google alone

# Outline

- Today we'll cover speed and scaling, including:
  - **Crawler design**
  - Inverted-index construction
  - Distributed search architecture
  - Deduplication

# Crawlers

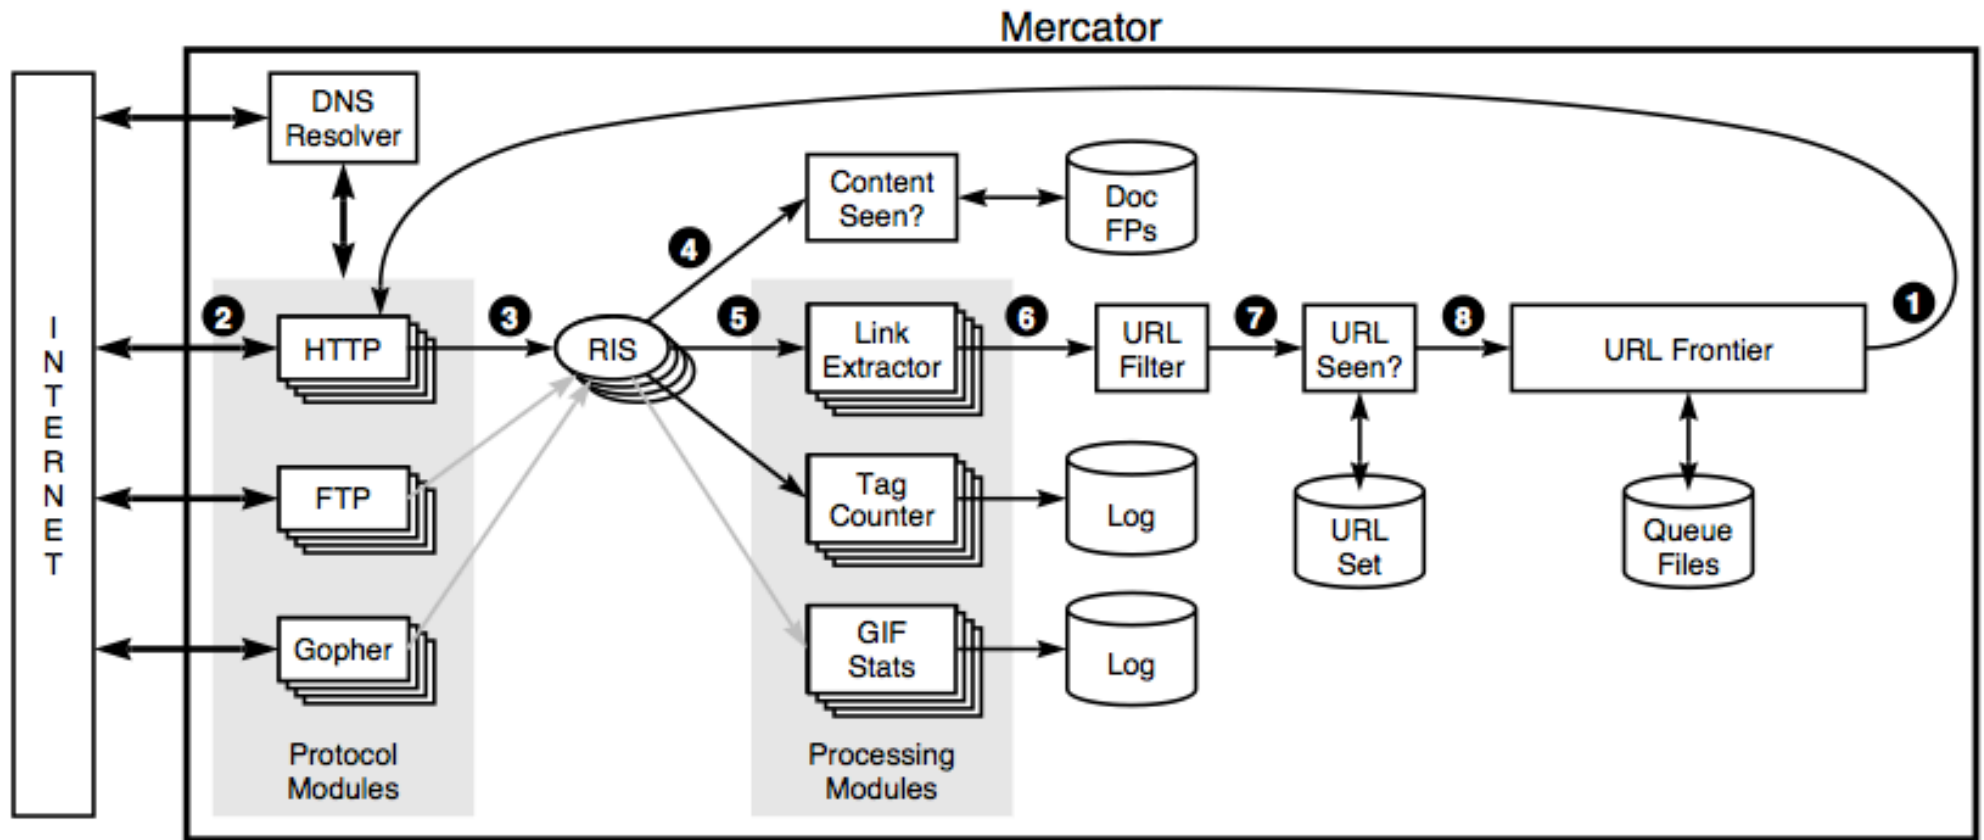
- To build a graph of the web, we need a program that visits every web page
- Web Robots AKA Web Wanderers, Crawlers, or Spiders
- Example: `googlebot`
  - <http://www.robotstxt.org/db/googlebot.html>

# Crawler Design

- To build a graph of the web, we need a program that visits every web page
- Web Robots AKA Web Wanderers, Crawlers, or Spiders
- Discussion: how would you do this?
- What data structures and algorithms would you use?

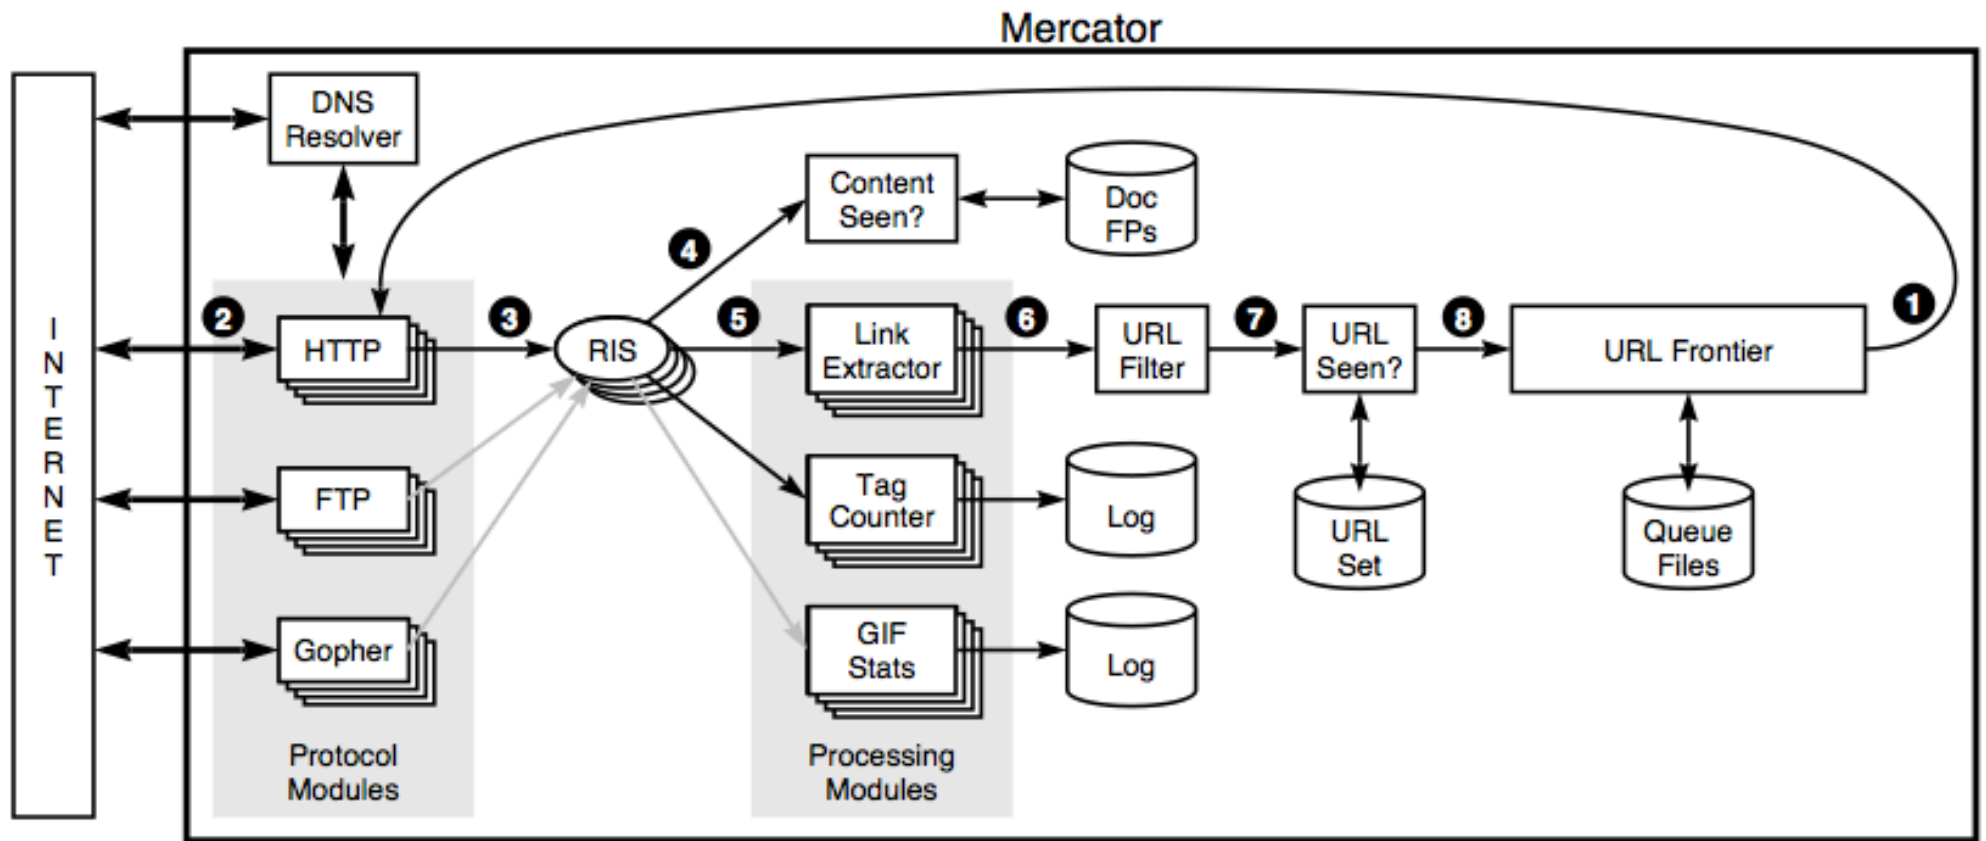
# Mercator

- Web crawler example: Mercator
- Mercator was the AltaVista crawler (1998)
- Exceptionally well-documented and useful, despite the many years
- Starts with seed URLs



1. Remove URL from queue
2. Network protocols
3. Read w/ RewindInputStream (RIS)
4. Has document been seen before?





5. Extract links
6. Download new URL?
7. Has URL been seen before?
8. Add URL to frontier
- GOTO 1

# What is Crawled?

- All static web pages
  - Including selected non-HTML pages
  - Unless restricted by robots.txt
- What about dynamic web pages?
  - Crawler visits, but can index properly only if it can understand what the page is about

# User agent

- When a browser or robot visits a page, it identifies itself with a `User-agent` string
  - For example, check yours out:
  - <http://www.whatismyuseragent.net/>
- Example from Google Chrome:
  - `Mozilla/5.0 (Macintosh; Intel Mac OS X 10_11_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/48.0.2564.103 Safari/537.36`
  - Previously used to indicate compatibility with the Mozilla rendering engine
  - During the "browser wars", some web sites would only send advanced features to some user agents

# User agent

- You can spoof your user agent 😊
- `curl -A "Mozilla/5.0"`  
`http://www.whatismyuseragent.net/`
- User agent switcher plug in for Chrome



# User agent

- Similar to a browser, when robot visits a page, it identifies itself with a `User-agent`
  - Just like Firefox, Chrome, Safari, etc.
- You can request that robots not visit your site with `/robots.txt`

# /robots.txt

User-agent: \*

Disallow: /

- User-agent: \*
- means this section applies to all robots.
- Disallow: /
- tells the robot that it should not visit any pages on the site.

# /robots.txt

User-agent: Googlebot-Image

Disallow: /

- Tell Google Image search not to include images from your website

# /robots.txt

User-agent:googlebot

Disallow:

User-agent: \*

Disallow:/private/

Disallow:/~jag/pvt.html

- Default is to allow
- More specific Disallow applies



# /robots.txt

- robots can ignore your /robots.txt
  - Malware robots that scan the web for security vulnerabilities
  - Email address harvesters used by spammers
- /robots.txt file is a publicly available file
  - Anyone can see what sections of your server you don't want robots to use
- /robots.txt directives can't prevent references to your URLs from other sites
  - A robot could navigate directly to a page from another website

# Outline

- Today we'll cover speed and scaling, including:
  - Crawler design
  - **Inverted-index construction**
  - Distributed search architecture
  - Deduplication

# Serving Results - speed

- After crawling is finished, we have a (big) database of documents
- To serve a search request, we need the terms in each doc
- You could just run `grep`
  - What is the complexity?
- How could we make this faster?

# Serving Results - speed




- After crawling is finished, we have a (big) database of documents
- To server a search request, we need the terms in each doc
- You could just run `grep`
  - What is the complexity?
  - $O(N)$ , where  $N$  is the total size of all web docs!
- How could we make this faster?

# Inverted Index

- A forward index is a list of words in each document
  - Doc -> words
- An *inverted index* maps words to docs that contain those words
  - Word -> docs
- Basic idea is: for each word, list all the documents where that word can be found
- Key to fast query processing

# Inverted Index

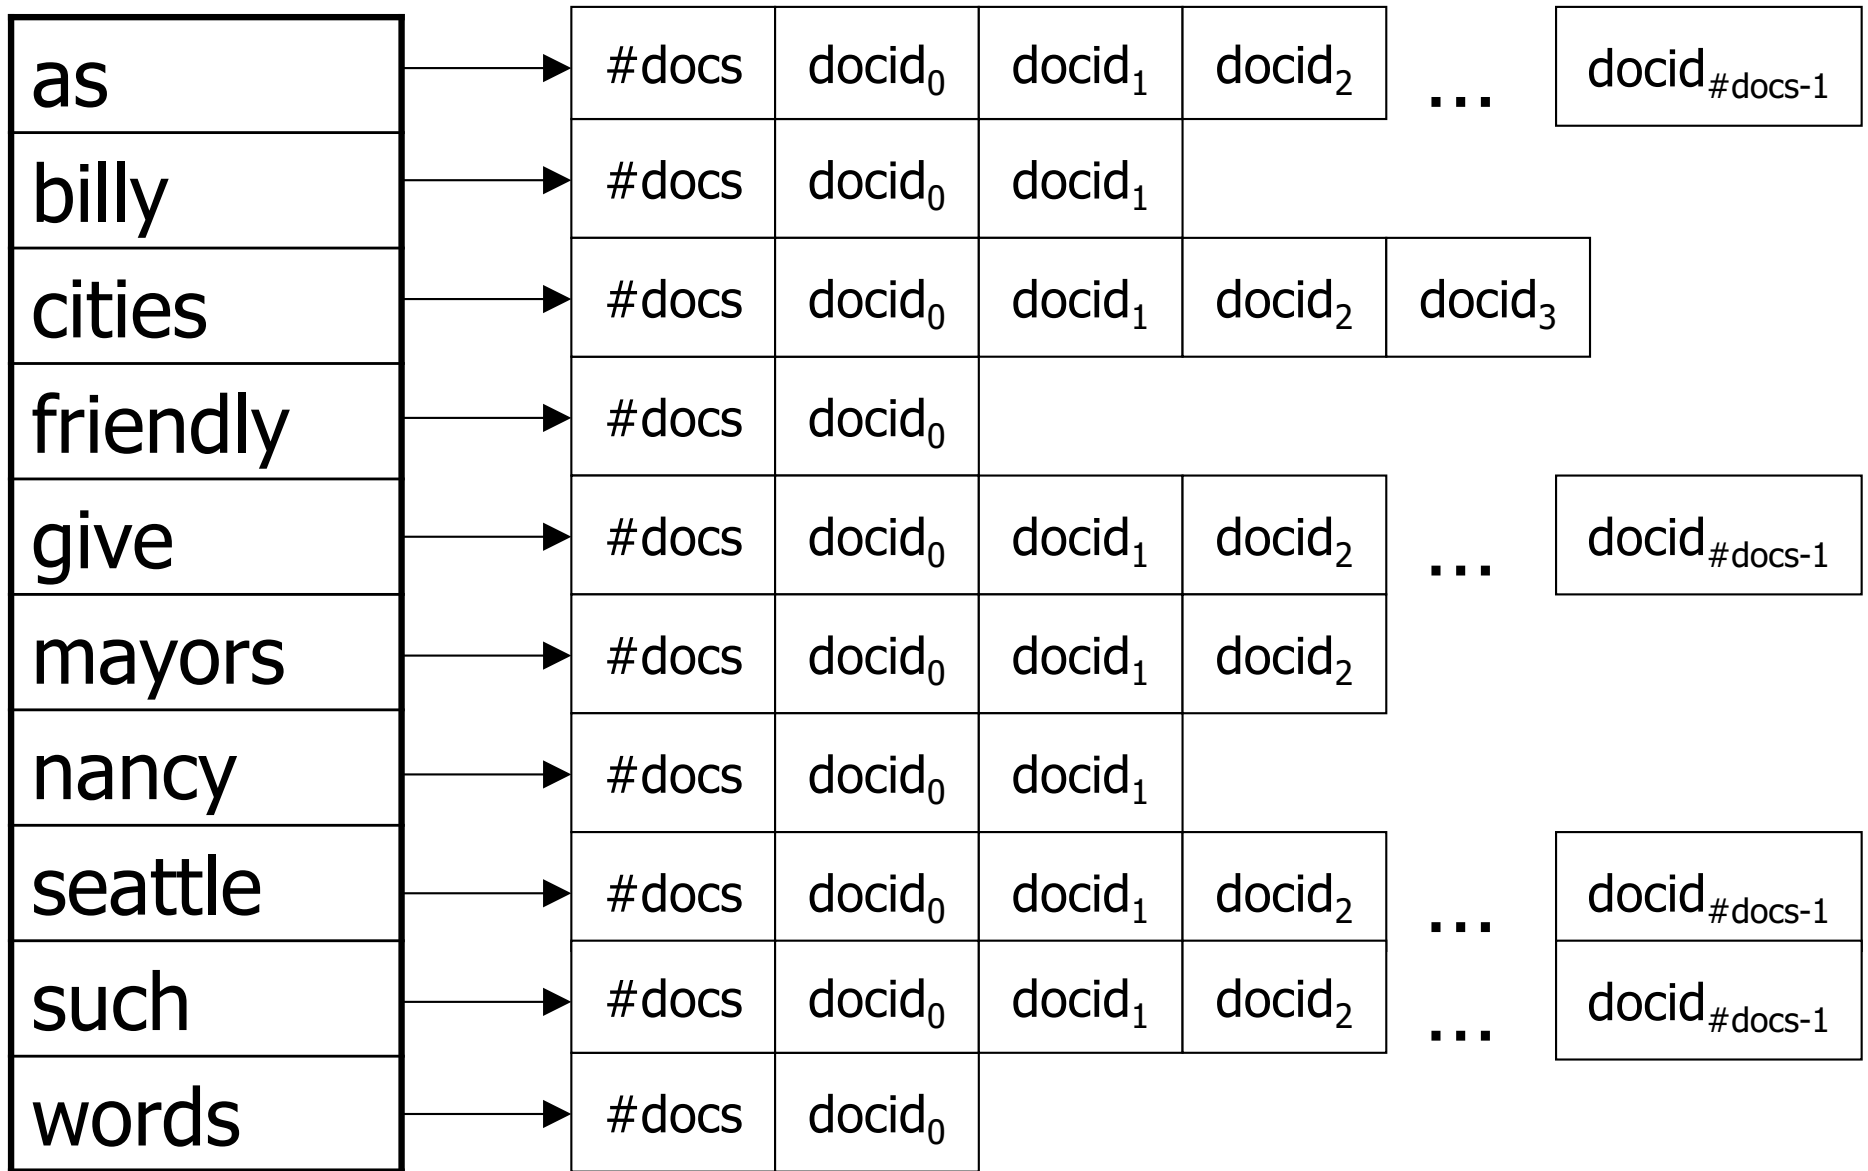
- Inverted indexes were around before computers
- Example: concordance
- List of every word, in alphabetical order
- Constructed manually before computers!!!



# Inverted Index

- openshakespeare.org concordance: horatio

#	Work	Character	Line	Text
1	<b>Hamlet</b> [I, 1]	<b>Bernardo</b>	13	Well, good night. If you do meet <b>Horatio</b> and Marcellus, The rivals of my watch, bid them make haste.
2	<b>Hamlet</b> [I, 1]	<b>(stage directions)</b>	16	Enter <b>Horatio</b> and Marcellus.
3	<b>Hamlet</b> [I, 1]	<b>Bernardo</b>	26	Say- What, is <b>Horatio</b> there ?
4	<b>Hamlet</b> [I, 1]	<b>Bernardo</b>	29	Welcome, <b>Horatio</b> . Welcome, good Marcellus.
5	<b>Hamlet</b> [I, 1]	<b>Marcellus</b>	32	<b>Horatio</b> says 'tis but our fantasy, And will not let belief take hold of him Touching this dreaded sight, twice seen of us. Therefore I have entreated him along, With us to watch the minutes of this night, That, if again this apparition come, He may approve our eyes and speak to it.
6	<b>Hamlet</b> [I, 1]	<b>Marcellus</b>	54	Thou art a scholar; speak to it, <b>Horatio</b> .





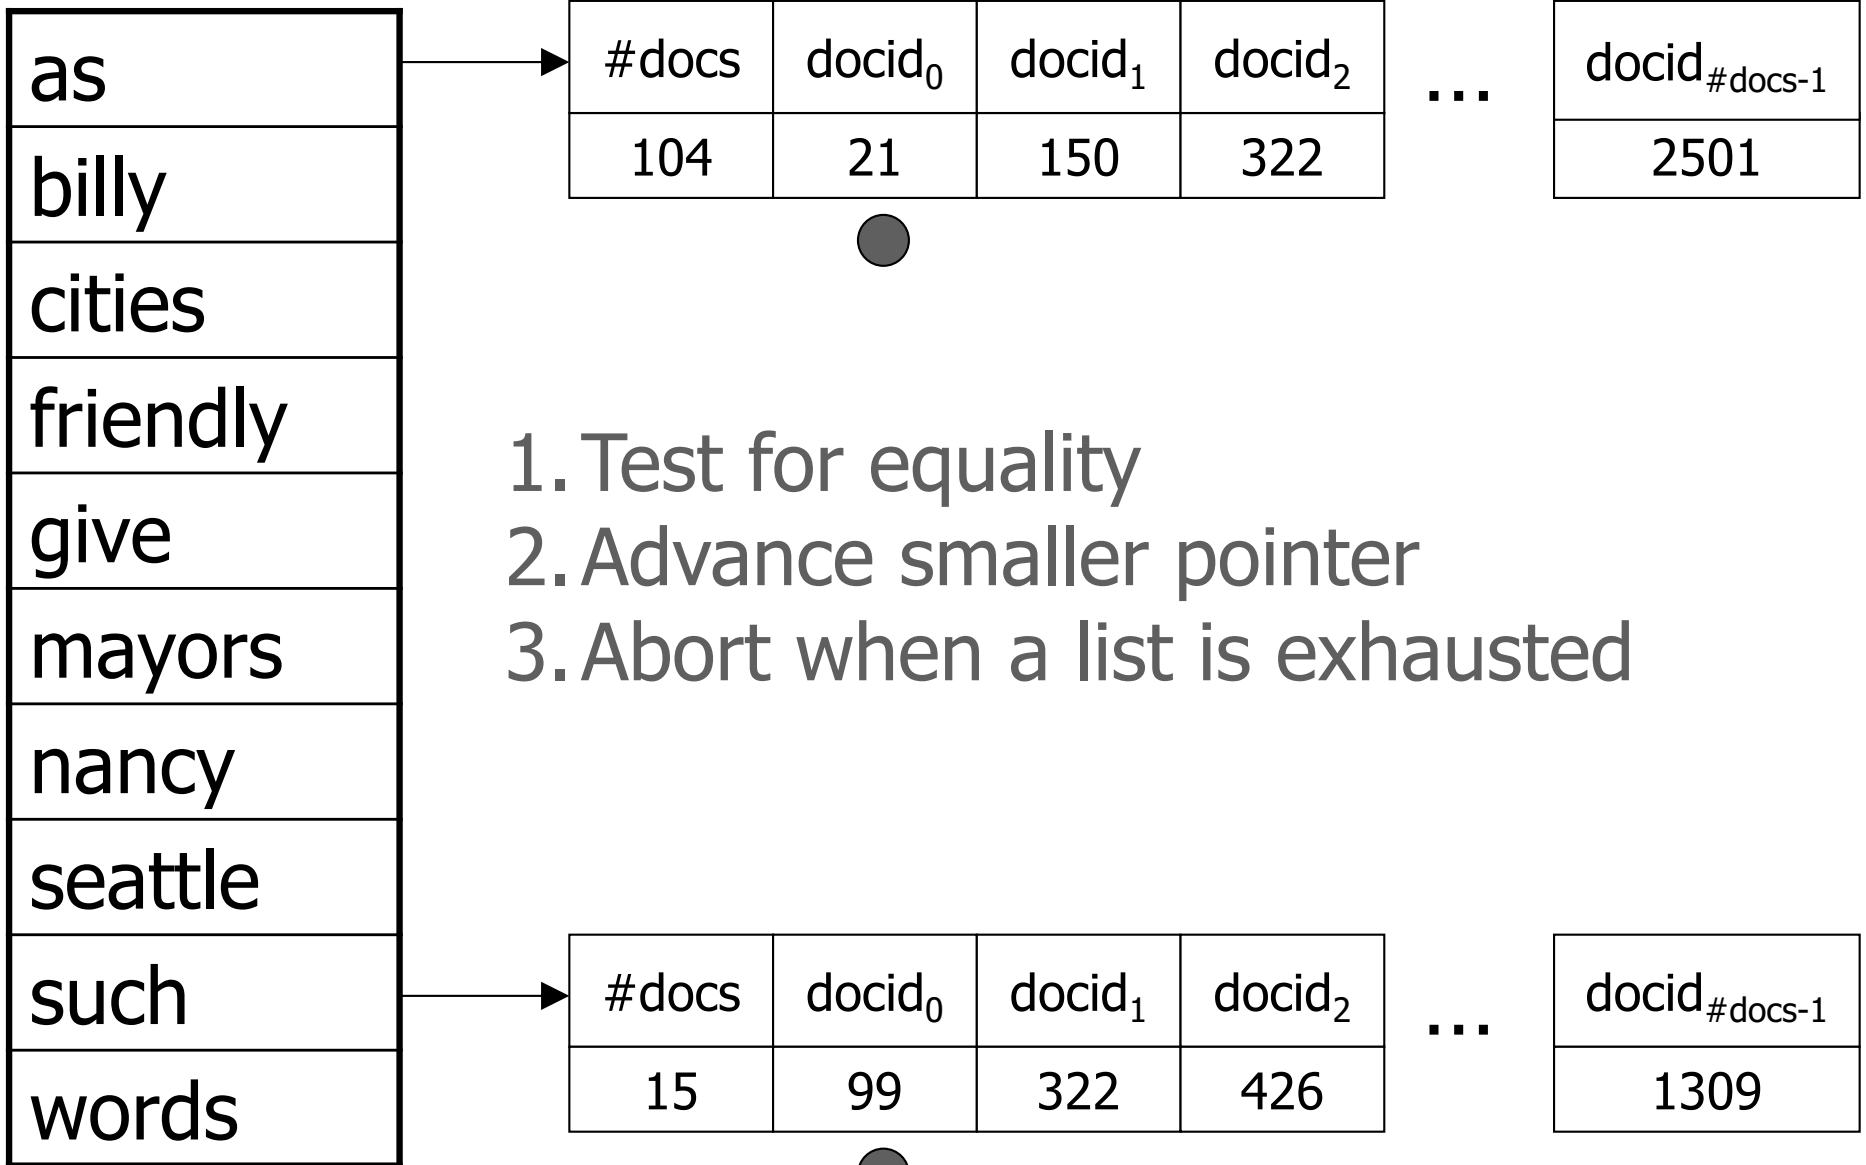
# Exercise

- Suggest an algorithm for finding a list of docs for a search
- Example: "such as"
  - Looking for keywords, not phrase
- HINT: does it help if the docs are sorted by ID?

# Exercise

- Suggest an algorithm for finding a list of docs for a search
- Example: "such as"
- One solution: merge intersection

Query: **such as**



**Returned docs:** 322

# Inverted Index Construction

- Inverted index is very large.
  - Full text index usually larger than the text.
  - How can we build it efficiently?
- Remember:
  - Disk seeks are very expensive (5ms)
  - Continuous disk reads or writes are OK (50-120MB/sec)
  - Machines can have a lot of memory (often hundreds of GB), but disk is always much cheaper
  - Input is the tokenized document set

# Basic Tasks

1. Compile term-termid, doc-docid maps
  2. Assemble all termid-docid pairs
  3. Sort pairs first by termid, then docid
  4. Write out in inverted-index form
- **EASY!**
  - Well, not if docs won't fit into memory

# Block sort-based indexing

- *External sort algorithms* work on sets larger than memory

- Block-Sort-Based Index Algorithm:

$n = 0$

While **docsRemain**

**n++**

**block** = ParseNextBlock()

BSBI-Invert(**block**)

WriteToDisk(block,  $f_n$ )

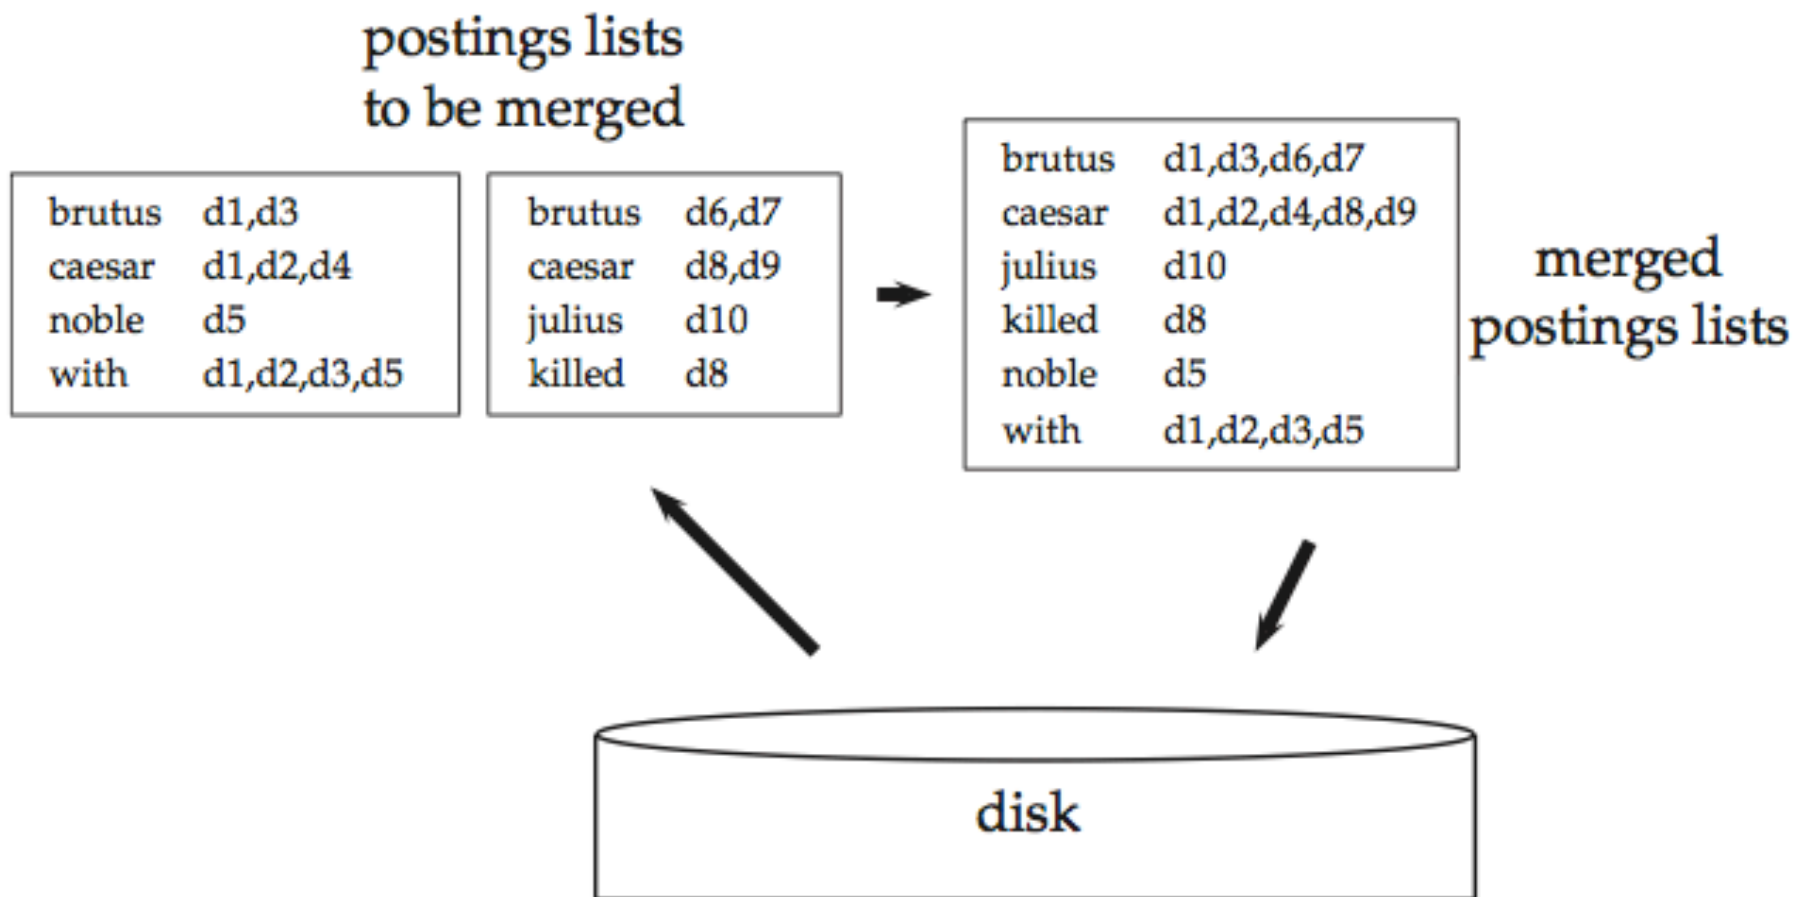
MergeBlocks( $f_1, \dots, f_n$ )  $\Rightarrow f_{\text{merged}}$

# Block sort-based indexing

- ParseNextBlock accumulates termid-docid pairs in memory until block is full
- BSBI-Invert generates small in-memory inverted index
- So: we build a series of small in-memory inverted indexes, writing each one to disk
- Finally: we merge them

# Block merging

Merging two lists:






# Outline

- Today we'll cover speed and scaling, including:
  - Crawler design
  - Inverted-index construction
  - **Distributed search architecture**
  - Deduplication


# Distributed Searching

- Not even the inverted index is small enough for one machine to handle it
  - Billions of docs
  - Hundreds of millions of queries
- Also, what if the machine fails?
- Need to parallelize query processing
  - Segment by document
  - Segment by search term

# Scaling - doc segmenting



# Scaling - term segmenting



# Segmentation

- What are the trade-offs of segment by documents vs. segment by term?
- What happens if a machine dies?

# Segmentation

- Segment by document
  - Easy to partition (just MOD the docid)
  - Easy to add new documents
  - If machine fails, quality goes down but queries don't die
- Segment by term
  - Harder to partition (terms uneven)
  - Trickier to add a new document (need to touch many machines)
  - If machine fails, search term might disappear, but not critical pages (e.g., [yahoo.com/index.html](http://yahoo.com/index.html))

# Outline

- Today we'll cover speed and scaling, including:
  - Crawler design
  - Inverted-index construction
  - Distributed search architecture
  - **Deduplication**

# Deduplication

- How can you be sure a web page is worth indexing?
  - Has it changed meaningfully?
  - A clone of another site? (Weirdly common)
- Two problems to solve:
  1. Are these two documents duplicates?
  2. Find all duplicates



# Are these two documents duplicates?

- Pages can have thousands of words. Comparing a pair of pages takes time.
  - How can you generate a page fingerprint?
  - What about a near-fingerprint?



# Shingling

- Compute the k-shingles for a page
- “I think EECS 485 is a great class”
  - For  $k=3$ : “I think EECS”, “think EECS 485”, “EECS 485 is”, etc
- If docs share lots shingles, they’re dups
- Document now a *set* of shingles
- Convert a document comparison problem into a set comparison problem
  - How similar are two sets?

# Jaccard similarity coefficient

- Jaccard similarity coefficient compares the similarity of the two sets of shingles (A and B)
- Size of the intersection / size of the union

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$



- 0 for disjoint sets, 1 for equal sets
- What is the complexity?
- Assume A and B are size  $O(N)$

# Jaccard similarity coefficient

- What is the complexity?
- Assume A and B are size O(N)

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}.$$

- O(N) time with O(N) space using hash tables

# Jaccard similarity coefficient

- Computing the Jaccard similarity coefficient is too slow for large documents
- Let's estimate it
- First, a question:
- Pick a random shingle from  $A \cup B$
- What is the probability that this shingle is in the intersection?

# Jaccard similarity coefficient

- Pick a random shingle from  $A \cup B$
- What is the probability that this shingle is in the intersection?

$$\frac{|A \cap B|}{|A \cup B|}$$

- That's the Jaccard similarity coefficient

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

# Selecting shingles

- How can we efficiently select a random shingle that is present in at least one of A or B?
  - In other words, a random shingle from  $A \cup B$
- A similar problem: select a random element from an array
  - Size is not known
  - Hint: let's say you have a perfect hash function

# Min hash

- Hash each element and choose the element that corresponds to the min of the hashed values
  - The hash function maps inputs uniformly over the output
  - Selecting the  $\min(h(x))$  is the same as selecting a random item  $x$ !
- This is called min hash



# Min hash and duplicate detection

- Back to doc A and doc B
- Want to compare a random shingle that is present in at least one doc (A or B)
- Hash all the shingles
  - $h(\text{shingle}) \rightarrow \text{integer}$
- $\min(h(\text{shingle}))$  corresponds to a random shingle
- This shingle will be in A and B with probability

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}.$$

- Again, the Jaccard similarity coefficient

# Min Hash Idea

- Now, pick the first  $k$  hashed values
- Alternatively, use  $k$  hash functions, and pick the min of each
- Finally, set union with  $k$  values instead of  $N$  values

# Information Retrieval

- At the heart of web search
- Document-vector model, network model
- Metrics: precision, recall, Kendall's Tau
- Implementation:
  - Crawler
  - Inverted Index
  - Deduplication: minhash, Jacard Coefficient.
- Basis for project 5
- Semester-long course in EECS 486